

Our goal is an optimisation of operations of type
 $g(f(...), ...)$

where f is a function that produces a big matrix
and g is a reduce function that makes the
big matrix small.

Example 1:

$$\boxed{A^T \times B \times v}, \text{ where } A, B \in \mathbb{R}^{M \times N} \text{ and } v \in \mathbb{R}^N$$

so our f here is

$$f: \mathbb{R}^{N \times M} \times \mathbb{R}^{M \times N} \rightarrow \mathbb{R}^{N \times N}$$
$$A^T, B \mapsto A^T \times B$$

and

$$g: \mathbb{R}^{N \times N} \times \mathbb{R}^N \rightarrow \mathbb{R}^N$$
$$(A^T \times B), v \mapsto (A^T \times B) \times v$$

the problematic matrix is a result of $f(A, B)$

Example 2: (Distance)

$$\boxed{f(a, b) = C \text{ and } g(C, v) = C \times v}$$

where $C \in \mathbb{R}^{N \times N}$ and $C_{ij} = [a_i - b_j]_{ij}$,
 $a \in \mathbb{R}^N, b \in \mathbb{R}^N, v \in \mathbb{R}^N$

Example 3: (multiplication)

The same as in Example 2, but

$$C_{ij} = [a_i * b_j]_{ij}$$

Example 4: (Squared distance)

$$f(A, B) = C, \text{ where } A \in \mathbb{R}^{N \times M}, B \in \mathbb{R}^{N \times M}$$

$$C_{ij} = \left[\left(\sum_k A_{ik} - B_{jk} \right)^2 \right]_{ij}, \text{ s.t. } C \in \mathbb{R}^{N \times N}$$

$$g(C, v) = C \times v$$

In Example 1, for $A^T \times B \times v$ we could simply re-arrange computation, because of associativity property of the matrix product, so instead of computing $(A^T \times B)$ first that gives $N \times N$ matrix, we could compute $(B \times v)$ first that gives N vector, and then $A^T \times (B \times v)$.


It reduces the memory requirements and computational cost.

Re-arranging is a partial solution, the generic solution would be chunking / splitting / partitioning big matrix computations and never materialise them at once on the same device. This could be referred as cache-wise computation.

E.g. for matrix-vector multiplication (MVM):

$$f(\dots) \in \mathbb{R}^{N \times N}$$

Diagram 1



$$= \text{vertical stack axis=0} \left(\begin{pmatrix} \boxed{1} \times \begin{pmatrix} 1 \\ 1 \end{pmatrix} \\ \boxed{3} \times \begin{pmatrix} 1 \\ 1 \end{pmatrix} \end{pmatrix} \oplus \begin{pmatrix} \boxed{2} \times \begin{pmatrix} 1 \\ 2 \end{pmatrix} \\ \boxed{4} \times \begin{pmatrix} 1 \\ 2 \end{pmatrix} \end{pmatrix} \right) \in \mathbb{R}^N$$

\oplus - is an element-wise addition operation

Now let's imagine that matrix $f(\dots)$ is a result of outer product (xy^T) where $x \in \mathbb{R}^N$ and $y \in \mathbb{R}^N$, and $xy^T \in \mathbb{R}^{N \times N}$

So for $f(x, y) = xy^T$:




Diagram 1 becomes :

$$= \text{vertical stack axis} = 0 \left(\begin{array}{c} \left(\begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \times \begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \times \begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \end{array} \right) \oplus \left(\begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \times \begin{array}{c} 2 \\ \vdots \\ 2 \end{array} \times \begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \end{array} \right) \\ \hline \left(\begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \times \begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \times \begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \end{array} \right) \oplus \left(\begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \times \begin{array}{c} 2 \\ \vdots \\ 2 \end{array} \times \begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \end{array} \right) \end{array} \right) \in \mathbb{R}^N$$

$$\left(\begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \times \begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \times \begin{array}{c} 1 \\ \vdots \\ 1 \end{array} \right)$$

This is a super efficient computation as 2nd dot product gives a scalar and 1st dot product becomes simply an element-wise multiplication of a scalar and a vector.

The total cost of operation is $O(N)$ and memory $O(N)$